## Drone-Based Object Detection and Tracking: Implementation and Analysis Using the Visdrone Dataset by Yolov8

Student: AmirHossein Eslami

Iran University Of Science And Technology



## Spring 2024

## **Table of Contents**

- 1. Introduction to the Visdrone Dataset
- 2. Structure of the Visdrone Dataset
- 3. Applications
- 4. Dataset YAML
- 5. Information from the Visdrone Dataset GitHub Page
- 6. Setting up and Installing Ultralytics
- 7. Training the Model in Google Colab
- 8. Displaying in TensorBoard
- 9. Model Prediction Analysis

The dataset we need to use to detect vehicles from above is one that consists of aerial images. Among the ready-to-use datasets available on the Ultralytics website, the Visdrone dataset meets our requirements.

Introduction to the Visdrone Dataset To get more familiar with this dataset, we visit the Ultralytics website and open the relevant dataset section. "The VisDrone dataset is a large-scale benchmark created by the AISKYEYE team at the Machine Learning and Data Mining Lab, Tianjin University, China. It includes accurately annotated ground-truth data for various computer vision tasks related to drone-based image and video analysis.

```
# Train/val/test sets as 1) dir: path/to/imgs, 2) file: path/to/imgs.txt, or 3) list: [path/to/imgs1, path/to/imgs2, ..]
path: ../datasets/VisDrone # dataset root dir
train: VisDrone2019-DET-train/images # train images (relative to 'path') 6471 images
val: VisDrone2019-DET-val/images # val images (relative to 'path') 548 images
test: VisDrone2019-DET-test-dev/images # test images (optional) 1610 images
```

```
# Classes
names:
0: pedestrian
1: people
2: bicycle
3: car
4: van
5: truck
6: tricycle
7: awning-tricycle
8: bus
9: motor
```

VisDrone consists of 288 video clips with 261,908 frames and 10,209 still images captured by various cameras mounted on drones. The dataset covers a wide range of aspects, including location (14 different cities across China), environment (urban and rural), objects (pedestrians, vehicles, bicycles, etc.), and density (sparse and crowded scenes). The dataset was collected using different drone platforms under various scenarios and weather and lighting conditions. These frames are manually annotated with over 2.6 million bounding boxes for objects like pedestrians, cars, bicycles, and tricycles. Features like scene visibility, object class, and occlusion are also provided for better data utilization."

Structure of the Visdrone Dataset The VisDrone dataset is organized into five main subsets, each focusing on a specific task:

- 1. Task 1: Object Detection in Images
- 2. Task 2: Object Detection in Videos
- 3. Task 3: Single Object Tracking
- 4. Task 4: Multiple Object Tracking
- 5. Task 5: Crowd Counting

Applications The VisDrone dataset is widely used for training and evaluating deep learning models in drone-based computer vision tasks such as object detection, object tracking, and crowd counting. The diverse sensor data, annotations, and dataset features make it a valuable resource for researchers and professionals in drone-based computer vision.

Dataset YAML A YAML (Yet Another Markup Language) file is used to define the dataset configuration. It includes information about dataset paths, classes, and other relevant details. For the Visdrone dataset, the VisDrone.yaml file is maintained at this link, which we will review next.

Information from the Visdrone Dataset GitHub Page This section provides the basic information needed to access and use the dataset, such as the dataset address and images used for training. Additionally, approximately 1610 images are available for testing the dataset.

The list of dataset classes is shown in this image. As seen, this dataset includes 10 classes of objects and entities that it can track and detect. Further, the download codes for the dataset are provided for developers.

Setting up and Installing Ultralytics First, we open a notebook in Google Colab. Then, we write the following code to install Ultralytics:

```
%pip install ultralytics
import ultralytics
ultralytics.checks()
```

After running the code, we observe that the installation is successful.

```
Ultralytics YOLOv8.1.43 🖋 Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4, 15102MiB)
Setup complete ☑ (2 CPUs, 12.7 GB RAM, 28.9/78.2 GB disk)
```

Next, we proceed to set up the required dataset.

Training the Model in Google Colab The number of chosen epochs is very high and time-consuming, so I reduced it by a factor of ten. Finally, the response we received indicates that the model has been trained:

```
10 epochs completed in 1.002 hours.

Optimizer stripped from runs/detect/train3/weights/last.pt, 87.6MB

Optimizer stripped from runs/detect/train3/weights/best.pt, 87.6MB

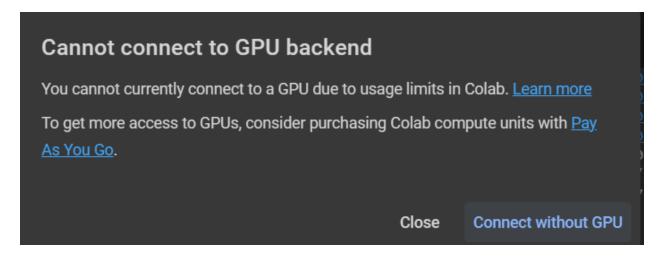
Validating runs/detect/train3/weights/best.pt...

Ultralytics YOLOv8.1.43  Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4, 15102MiB)

Model summary (fused): 268 layers, 43614318 parameters, 0 gradients, 164.9 GFLOPs

Class Images Instances Box(P R mAP50 mAP50-95): 22% 4/18 [00:37<02:34, 11.04s/it]
```

Unfortunately, due to extensive attempts to train, the free quota eventually ran out.

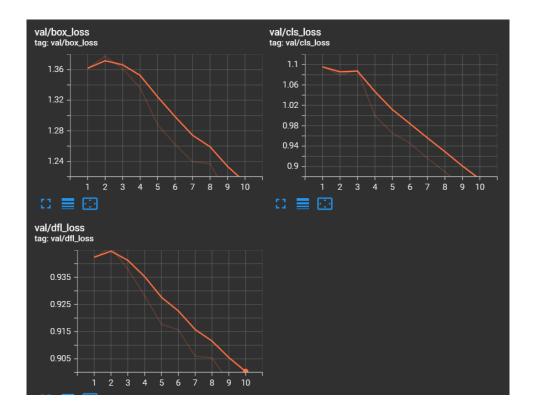


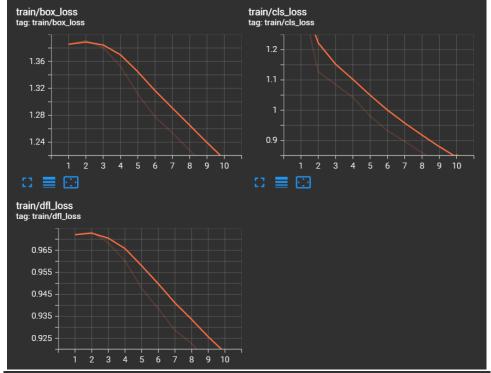
As a result, we have to perform all the steps again on another account to continue.

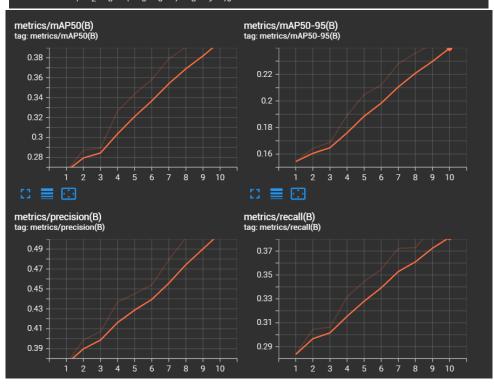
## Displaying in TensorBoard

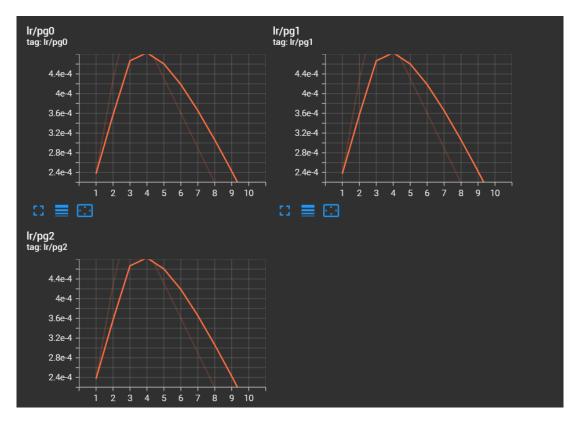
Model Prediction Analysis Well, the model is fully trained now, and we are ready to predict. All we need is an aerial video to use as a sample for prediction. Using the following command, we can perform the prediction:

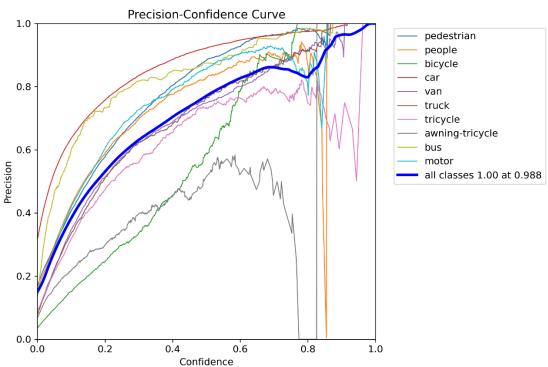
As shown, the prediction has started, and given the large number of video frames, it will take a long time to process each frame and perform the prediction.

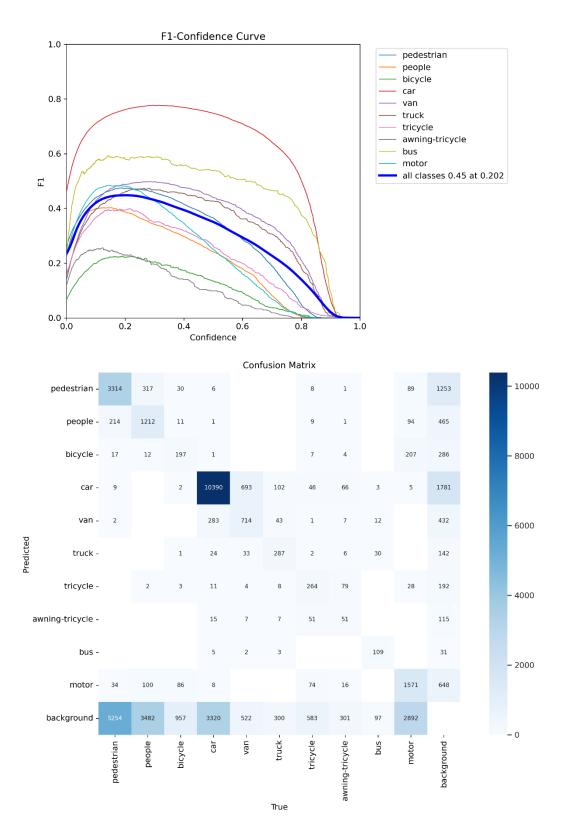


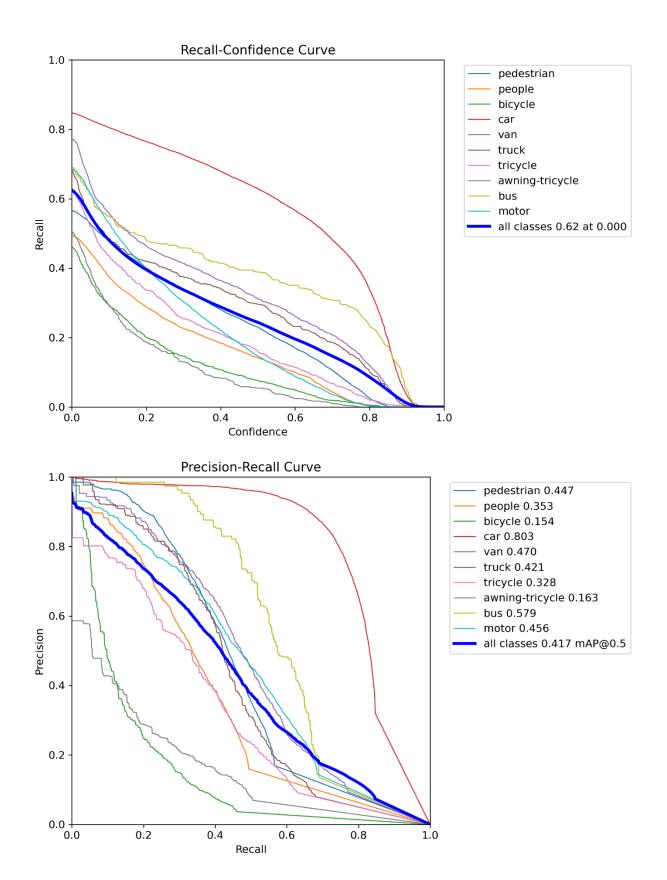












```
0: 384x640 1 pedestrian, 24 cars, 1 van, 4 trucks, 39.6ms
0: 384x640 1 pedestrian, 24 cars, 1 van, 3 trucks, 39.5ms
0: 384x640 1 pedestrian, 26 cars, 1 van, 3 trucks, 39.5ms
0: 384x640 2 pedestrians, 31 cars, 1 van, 6 trucks, 39.5ms
0: 384x640 2 pedestrians, 31 cars, 1 van, 3 trucks, 39.5ms
0: 384x640 2 pedestrians, 29 cars, 1 van, 3 trucks, 39.5ms
0: 384x640 1 pedestrian, 29 cars, 1 van, 4 trucks, 48.9ms
0: 384x640 1 pedestrian, 31 cars, 1 van, 3 trucks, 39.5ms
```

Finally, the prediction result is obtained, and it is observed that the prediction has been done very well, indicating that our model has been well-trained for object detection.

Here is part of the image obtained as a prediction result:

